

How many pods exist on the system? 0

```
controlplane ~ → kubectl get pods
No resources found in default namespace.
```

Create a new pod with the nginx image.

```
controlplane ~ → kubectl run nginx --image=nginx
pod/nginx created
```

How many pods are created now? 4

```
controlplane ~ → kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx         1/1     Running   0           36s
newpods-m5zn8 1/1     Running   0           16s
newpods-qdxrn 1/1     Running   0           16s
newpods-vszzv 1/1     Running   0           16s
```

What is the image used to create the new pods? busybox

```
controlplane ~ → kubectl describe pod newpods-m5zn8
Name:          newpods-m5zn8
Namespace:     default
Priority:       0
Service Account: default
Node:          controlplane/172.25.0.68
Start Time:    Tue, 25 Apr 2023 03:42:01 +0000
Labels:        tier=busybox
Annotations:    <none>
Status:        Running
IP:            10.42.0.12
IPs:
  IP:          10.42.0.12
Controlled By: ReplicaSet/newpods
Containers:
  busybox:
    Container ID:  containerd://460350097c800266eabeabc9fc8db6cc164174aca6f34bd79fc061
    Image:         busybox
    Image ID:      docker.io/library/busybox@sha256:b5d6fe0712636ceb7430189de28819e15
    Port:         <none>
    Host Port:    <none>
    Command:
      sleep
      1000
    State:        Running
      Started:    Tue, 25 Apr 2023 03:42:03 +0000
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-xlb4s (ro)
Conditions:
  Type            Status
  Initialized     True
  Ready           True
  ContainersReady True
  PodScheduled    True
Volumes:
```

Which nodes are these pods placed on? Controlplane

Run the command `kubectl describe pod newpods-<id>` and look at the node field.

Alternatively run `kubectl get pods -o wide` and check for the node the pod is placed on.

```
controlplane ~ → kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
nginx	1/1	Running	0	3m2s	10.42.0.9	controlplane	<none>		<none>	
newpods-m5zn8	1/1	Running	0	2m42s	10.42.0.12	controlplane	<none>		<none>	
newpods-qdxrn	1/1	Running	0	2m42s	10.42.0.11	controlplane	<none>		<none>	
newpods-vszzv9	1/1	Running	0	2m42s	10.42.0.10	controlplane	<none>		<none>	

How many containers are part of the pod webapp? 2

```
controlplane ~ → kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
nginx	1/1	Running	0	3m47s	10.42.0.9	controlplane	<none>		<none>	
newpods-m5zn8	1/1	Running	0	3m27s	10.42.0.12	controlplane	<none>		<none>	
newpods-qdxrn	1/1	Running	0	3m27s	10.42.0.11	controlplane	<none>		<none>	
newpods-vszzv9	1/1	Running	0	3m27s	10.42.0.10	controlplane	<none>		<none>	
webapp	1/2	ImagePullBackOff	0	27s	10.42.0.13	controlplane	<none>		<none>	

What images are used in the new webapp pod? Nginx & agentx

```
controlplane ~ → kubectl describe pods webapp
```

Name: webapp
Namespace: default
Priority: 0
Service Account: default
Node: controlplane/172.25.0.68
Start Time: Tue, 25 Apr 2023 03:45:01 +0000
Labels: <none>
Annotations: <none>
Status: Pending
IP: 10.42.0.13
IPs:
 IP: 10.42.0.13
Containers:
 nginx:
 Container ID: containerd://518ed25e679f624a1bc6ba0ae5865fff7efaf
 Image: nginx
 Image ID: docker.io/library/nginx@sha256:63b44e8ddb83d5dd80
 Port: <none>
 Host Port: <none>
 State: Running
 Started: Tue, 25 Apr 2023 03:45:04 +0000
 Ready: True
 Restart Count: 0
 Environment: <none>
 Mounts:
 /var/run/secrets/kubernetes.io/serviceaccount from kube-api-acc
 agentx:
 Container ID:
 Image: agentx
 Image ID:
 Port: <none>
 Host Port: <none>
 State: Waiting
 Reason: ErrImagePull

What is the state of the container agentx in the pod webapp? Error or waiting

```
controlplane ~ → kubectl describe pods webapp
Name:          webapp
Namespace:     default
Priority:       0
Service Account: default
Node:          controlplane/172.25.0.68
Start Time:    Tue, 25 Apr 2023 03:45:01 +0000
Labels:        <none>
Annotations:   <none>
Status:        Pending
IP:            10.42.0.13
IPs:
  IP: 10.42.0.13
Containers:
  nginx:
    Container ID:  containerd://518ed25e679f624a1bc6ba0ae5865fff7efa
    Image:         nginx
    Image ID:      docker.io/library/nginx@sha256:63b44e8ddb83d5dd80
    Port:         <none>
    Host Port:    <none>
    State:        Running
      Started:    Tue, 25 Apr 2023 03:45:04 +0000
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-acc
  agentx:
    Container ID:
    Image:         agentx
    Image ID:
    Port:         <none>
    Host Port:    <none>
    State:        Waiting
      Reason:     ErrImagePull
    Ready:        False
    Restart Count: 0
    Environment:  <none>
```

Why do you think the container agentx in pod webapp is in error? Docker image doesn't exist

Try to figure it out from the events section of the pod.

Run the command `kubectl describe pod webapp` and look under the events section.

An image by that name does not exist in DockerHub.

```
node.kubernetes.io/unreacheable: Node exists for 300s
Events:
-----
Type      Reason      Age      From      Message
-----
Normal    Scheduled   64s      default-scheduler    Successfully assigned default/webapp to controlplane
Normal    Pulling     63s      kubelet    Pulling image "nginx"
Normal    Pulled      61s      kubelet    Successfully pulled image "nginx" in 1.852069473s (1.852087456s including waiting)
Normal    Created     61s      kubelet    Created container nginx
Normal    Started     61s      kubelet    Started container nginx
Normal    Pulling     21s (x3 over 61s)  kubelet    Pulling image "agentx"
Warning   Failed      21s (x3 over 60s)  kubelet    Failed to pull image "agentx": rpc error: code = Unknown desc = failed to pull and unpack image
agentx:latest: failed to resolve reference "docker.io/library/agentx:latest": pull access denied, repository does not exist or may require authorization: insufficient_scope: authorization failed
Warning   Failed      21s (x3 over 60s)  kubelet    Error: ErrImagePull
Normal    BackOff     7s (x3 over 59s)   kubelet    Back-off pulling image "agentx"
Warning   Failed      7s (x3 over 59s)   kubelet    Error: ImagePullBackOff
```

What does the READY column in the output of the `kubectl get pods` command indicate?
Running container / total container in pod

```
controlplane ~ → kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx         1/1     Running   0           8m36s
newpods-m5zn8 1/1     Running   0           8m16s
newpods-qdxrn 1/1     Running   0           8m16s
newpods-vszzv9 1/1     Running   0           8m16s
webapp        1/2     ImagePullBackOff 0           5m16s
```

Delete the webapp Pod.
Once deleted, wait for the pod to fully terminate

```
controlplane ~ ✗ kubectl delete pod webapp
pod "webapp" deleted
```

Create a new pod with the name redis and with the image redis123.
Use a pod-definition YAML file. And yes the image name is wrong!
We use `kubectl run` command with `--dry-run=client -o yaml` option to create a manifest file :-
After that, using `kubectl create -f` command to create a resource from the manifest file :-
Verify the work by running `kubectl get` command :-

```
controlplane ~ → kubectl run redis --image=redis123 --dry-run=client
pod/redis created (dry run)

controlplane ~ → kubectl run redis --image=redis123 --dry-run=client -o yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: redis
  name: redis
spec:
  containers:
  - image: redis123
    name: redis
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}

controlplane ~ → kubectl run redis --image=redis123 --dry-run=client -o yaml > redis-def.yaml

controlplane ~ → kubectl create -f redis-def.yaml
pod/redis created

controlplane ~ → kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx         1/1     Running   0           15m
newpods-m5zn8 1/1     Running   0           15m
newpods-qdxrn 1/1     Running   0           15m
newpods-vszzv9 1/1     Running   0           15m
redis         0/1     ErrImagePull 0           7s
```

Now change the image on this pod to redis.

```
controlplane ~ ❌ cat redis-def.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: redis
    name: redis
spec:
  containers:
  - image: redis
    name: redis
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
  status: {}

controlplane ~ ➔ kubectl apply -f redis-def.yaml
Warning: resource pods/redis is missing the kubectrl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.
pod/redis configured
```

REPLICASETS

How many ReplicaSets exist on the system?

```
controlplane ~ ➔ kubectl get pods
No resources found in default namespace.

controlplane ~ ➔ kubectl get rs
No resources found in default namespace.
```

How many PODs are DESIRED in the new-replica-set?

```
controlplane ~ ➔ kubectl get rs
NAME                DESIRED   CURRENT   READY   AGE
new-replica-set     4         4         0       13s
```

What is the image used to create the pods in the new-replica-set?

```
controlplane ~ ➔ kubectl describe replicaset
Name:                new-replica-set
Namespace:           default
Selector:             name=busybox-pod
Labels:               <none>
Annotations:          <none>
Replicas:             4 current / 4 desired
Pods Status:         0 Running / 4 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:             name=busybox-pod
  Containers:
    busybox-container:
      Image:           busybox777
      Port:            <none>
      Host Port:       <none>
      Command:
        sh
        -c
        echo Hello Kubernetes! && sleep 3600
      Environment:     <none>
      Mounts:          <none>
      Volumes:         <none>
Events:
  Type      Reason            Age   From                      Message
  ----      -
  Normal    SuccessfulCreate   72s   replicaset-controller     Created pod: new-replica-set-l5x69
  Normal    SuccessfulCreate   72s   replicaset-controller     Created pod: new-replica-set-nnpvz
  Normal    SuccessfulCreate   72s   replicaset-controller     Created pod: new-replica-set-5xbv6
  Normal    SuccessfulCreate   72s   replicaset-controller     Created pod: new-replica-set-5256z
```


How many PODs are READY in the new-replica-set?

```
controlplane ~ → kubectl get rs
NAME                DESIRED   CURRENT   READY   AGE
new-replica-set     4         4         0       2m12s
```

Why do you think the PODs are not ready? Image not exist

Run the command: `kubectl describe pods` and look at under the Events section.

```
Events:
  Type      Reason      Age           From          Message
  ----      -
Normal     Scheduled   3m3s         default-scheduler   Successfully assigned default/new-replica-set-nnpvz to controlplane
Normal     Pulling     102s (x4 over 3m2s)  kubelet          Pulling image "busybox777"
Warning    Failed      101s (x4 over 3m2s)  kubelet          Failed to pull image "busybox777": rpc error: code = Unknown desc = failed to pull and unpack image "docker.io/library/busybox777:latest": failed to resolve reference "docker.io/library/busybox777:latest": pull access denied, repository does not exist or may require authorization
: server message: insufficient_scope: authorization failed
Warning    Failed      101s (x4 over 3m2s)  kubelet          Error: ErrImagePull
Warning    Failed      76s (x6 over 3m1s)   kubelet          Error: ImagePullBackOff
Normal     BackOff     63s (x7 over 3m1s)   kubelet          Back-off pulling image "busybox777"
```

Delete any one of the 4 PODs.

How many PODs exist now?

```
controlplane ~ → kubectl get pods
NAME                READY   STATUS              RESTARTS   AGE
new-replica-set-nnpvz 0/1     ImagePullBackOff    0          4m31s
new-replica-set-5xbv6 0/1     ImagePullBackOff    0          4m31s
new-replica-set-5256z 0/1     ImagePullBackOff    0          4m31s
new-replica-set-l5x69 0/1     ImagePullBackOff    0          4m31s

controlplane ~ → kubectl delete pod new-replica-set-nnpvz
pod "new-replica-set-nnpvz" deleted

controlplane ~ → kubectl get pods
NAME                READY   STATUS              RESTARTS   AGE
new-replica-set-5xbv6 0/1     ImagePullBackOff    0          4m56s
new-replica-set-5256z 0/1     ImagePullBackOff    0          4m56s
new-replica-set-l5x69 0/1     ImagePullBackOff    0          4m56s
new-replica-set-9svpb 0/1     ErrImagePull        0          11s
```

Why are there still 4 PODs, even after you deleted one?

Replicaset ensures desired number of pods always run .

Create a ReplicaSet using the replicaset-definition-1.yaml file located at `/root/`.

There is an issue with the file, so try to fix it.

Run the command: `kubectl explain replicaset | grep VERSION` and correct the `apiVersion` for ReplicaSet.

Then run the command: `kubectl create -f /root/replicaset-definition-1.yaml`

```
controlplane ~ → cat replicaset-definition-1.yaml
apiVersion: v1
kind: ReplicaSet
metadata:
  name: replicaset-1
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: nginx
          image: nginx
```

```
controlplane ~ → kubectl explain replicaset | grep VERSION
VERSION: apps/v1
```

```
controlplane ~ → kubectl explain rs
KIND:      ReplicaSet
VERSION:   apps/v1
```

DESCRIPTION:

ReplicaSet ensures that a specified number of pod replicas are running at any given time.

FIELDS:

apiVersion <string>
APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info:
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources>

kind <string>
Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info:
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#types-kinds>

metadata <Object>
If the Labels of a ReplicaSet are empty, they are defaulted to be the same as the Pod(s) that the ReplicaSet manages. Standard object's metadata. More info:
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#metadata>

spec <Object>
Spec defines the specification of the desired behavior of the ReplicaSet. More info:
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status>

status <Object>

```
controlplane ~ → cat replicaset-definition-1.yaml
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: replicaset-1
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: nginx
          image: nginx
```

```
controlplane ~ → kubectl create -f replicaset-definition-1.yaml
replicaset.apps/replicaset-1 created
```

Fix the issue in the replicaset-definition-2.yaml file and create a ReplicaSet using it.
This file is located at /root/.

```
controlplane ~ ✗ cat replicaset-definition-2.yaml
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: replicaset-2
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```



```
controlplane ~ → cat replicaset-definition-2.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: replicaset-2
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: nginx
  template:
    metadata:
      labels:
        tier: nginx
    spec:
      containers:
      - name: nginx
        image: nginx

controlplane ~ → kubectl apply -f replicaset-definition-2.yaml
replicaset.apps/replicaset-2 created
```

Delete the two newly created ReplicaSets - replicaset-1 and replicaset-2
Run the command: `kubectl delete replicaset <replicaset-name>` or `kubectl delete -f <file-name>.yaml`

```
controlplane ~ → kubectl delete rs replicaset-1 replicaset-2
replicaset.apps "replicaset-1" deleted
replicaset.apps "replicaset-2" deleted
```

Fix the original replica set new-replica-set to use the correct busybox image.
Either delete and recreate the ReplicaSet or Update the existing ReplicaSet and then delete all PODs, so new ones with the correct image will be created.

```
controlplane ~ → kubectl edit rs new-replica-set
replicaset.apps/new-replica-set edited
```

```
#
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  creationTimestamp: "2023-04-25T04:02:11Z"
  generation: 1
  name: new-replica-set
  namespace: default
  resourceVersion: "1070"
  uid: 2eda23f6-65c7-44b7-bdc2-938ad71d6f86
spec:
  replicas: 4
  selector:
    matchLabels:
      name: busybox-pod
  template:
    metadata:
      creationTimestamp: null
      labels:
        name: busybox-pod
    spec:
      containers:
      - command:
        - sh
        - -c
        - echo Hello Kubernetes! && sleep 3600
        image: busybox
        imagePullPolicy: Always
        name: busybox-container
        resources: {}
        terminationMessagePath: /dev/termination-log
        terminationMessagePolicy: File
      dnsPolicy: ClusterFirst
      restartPolicy: Always
```

:wq

```
controlplane ~ → kubectl delete pods new-replica-set-9svpb new-replica-set-15x69 new-replica-set-5256z new-replica-set-5xbv6
pod "new-replica-set-9svpb" deleted
pod "new-replica-set-15x69" deleted
pod "new-replica-set-5256z" deleted
pod "new-replica-set-5xbv6" deleted
```

```
controlplane ~ → kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
new-replica-set-pl9x2	1/1	Running	0	7s
new-replica-set-6xgbj	1/1	Running	0	8s
new-replica-set-p8s8t	1/1	Running	0	7s
new-replica-set-sfvzb	1/1	Running	0	8s

Scale the ReplicaSet to 5 PODs.

Use kubectl scale command or edit the replicaset using kubectl edit replicaset.

Run the command: kubectl edit replicaset new-replica-set, modify the replicas and then save the file OR run:

kubectl scale rs new-replica-set --replicas=5 to scale up to 5 PODs.

```
controlplane ~ → kubectl get rs
NAME                DESIRED    CURRENT    READY    AGE
new-replica-set     4          4          4        18m

controlplane ~ → kubectl scale rs new-replica-set --replicas=5
replicaset.apps/new-replica-set scaled
```

Now scale the ReplicaSet down to 2 PODs.

Use the kubectl scale command or edit the replicaset using kubectl edit replicaset.

Run the command: kubectl edit replicaset new-replica-set, modify the replicas and then save the file OR run: kubectl scale rs new-replica-set --replicas=2 to scale down to 2 PODs.

```
controlplane ~ → kubectl get rs
NAME                DESIRED    CURRENT    READY    AGE
new-replica-set     5          5          5        19m

controlplane ~ → kubectl scale rs new-replica-set --replicas=2
replicaset.apps/new-replica-set scaled

controlplane ~ → kubectl get pods
NAME                READY    STATUS    RESTARTS    AGE
new-replica-set-6xgbj 1/1      Running   0           2m48s
new-replica-set-sfvzb 1/1      Running   0           2m48s
new-replica-set-pl9x2 1/1      Terminating 0           2m47s
new-replica-set-p8s8t 1/1      Terminating 0           2m47s
new-replica-set-q9hwr 1/1      Terminating 0           88s
```

```
controlplane ~ → kubectl get pods
No resources found in default namespace.

controlplane ~ → kubectl get rs
No resources found in default namespace.

controlplane ~ → kubectl get deployments
No resources found in default namespace.

controlplane ~ → kubectl get deployments
NAME                READY    UP-TO-DATE    AVAILABLE    AGE
frontend-deployment 0/4      4              0            4s

controlplane ~ → kubectl get rs
NAME                DESIRED    CURRENT    READY    AGE
frontend-deployment-7fbf4f5cd9 4          4          0        14s

controlplane ~ → kubectl get pods
NAME                READY    STATUS    RESTARTS    AGE
frontend-deployment-7fbf4f5cd9-ff7cj 0/1      ImagePullBackOff 0           27s
frontend-deployment-7fbf4f5cd9-xqk9b 0/1      ImagePullBackOff 0           27s
frontend-deployment-7fbf4f5cd9-qfj84 0/1      ImagePullBackOff 0           27s
frontend-deployment-7fbf4f5cd9-kn2s9 0/1      ImagePullBackOff 0           27s
```

What is the image used to create the pods in the new deployment? Busybox888

```
controlplane ~ ❌ kubectl describe pods
Name:          frontend-deployment-7fbf4f5cd9-kn2s9
Namespace:     default
Priority:       0
Service Account: default
Node:          controlplane/172.25.0.26
Start Time:    Tue, 25 Apr 2023 04:41:51 +0000
Labels:        name=busybox-pod
               pod-template-hash=7fbf4f5cd9
Annotations:   <none>
Status:        Pending
IP:            10.42.0.12
IPs:
  IP:          10.42.0.12
Controlled By: ReplicaSet/frontend-deployment-7fbf4f5cd9
Containers:
  busybox-container:
    Container ID:
    Image:        busybox888
    Image ID:
    Port:         <none>
    Host Port:    <none>
    Command:
      sh
      -c
      echo Hello Kubernetes! && sleep 3600
    State:        Waiting
    Reason:       ImagePullBackOff
    Ready:        False
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-f85zd (ro)
```

Why do you think the deployment is not ready? Image not exist

```
node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
Type      Reason      Age          From          Message
----      -
Normal    Scheduled   89s         default-scheduler   Successfully assigned default/frontend-deployment-7fbf4f5cd9-qfj84 to controlplane
Normal    BackOff     17s (x4 over 87s) kubelet        Back-off pulling image "busybox888"
Warning   Failed      17s (x4 over 87s) kubelet        Error: ImagePullBackOff
Normal    Pulling     2s (x4 over 88s) kubelet        Pulling image "busybox888"
Warning   Failed      1s (x4 over 88s) kubelet        Failed to pull image "busybox888": rpc error: code = Unknown desc = failed to pull and unpack image "docker.io/library/busybox888:latest": failed to resolve reference "docker.io/library/busybox888:latest": pull access denied, repository does not exist or may require authorization: server message: insufficient_scope: authorization failed
Warning   Failed      1s (x4 over 88s) kubelet        Error: ErrImagePull
```

Create a new Deployment using the deployment-definition-1.yaml file located at /root/.
There is an issue with the file, so try to fix it.

```

controlplane ~ → kubectl explain deployments | head -n1
KIND:      Deployment

controlplane ~ → kubectl explain deployments
KIND:      Deployment
VERSION:   apps/v1

DESCRIPTION:
  Deployment enables declarative updates for Pods and ReplicaSets.

FIELDS:
  apiVersion    <string>
    APIVersion defines the versioned schema of this representation of an
    object. Servers should convert recognized schemas to the latest internal
    value, and may reject unrecognized values. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources

  kind <string>
    Kind is a string value representing the REST resource this object
    represents. Servers may infer this from the endpoint the client submits
    requests to. Cannot be updated. In CamelCase. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#types-kinds

  metadata      <Object>
    Standard object's metadata. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#metadata

  spec <Object>
    Specification of the desired behavior of the Deployment.

  status        <Object>
    Most recently observed status of the Deployment.

```

Create a new Deployment with the below attributes using your own deployment definition file.

Name: httpd-frontend;

Replicas: 3;

Image: httpd:2.4-alpine

```

controlplane ~ → cat dep-httpd.yaml
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: httpd-frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      name: httpd-frontend
  template:
    metadata:
      labels:
        name: httpd-frontend
    spec:
      containers:
        - name: httpd-frontend
          image: httpd:2.4-alpine

controlplane ~ → kubectl create -f dep-httpd.yaml
deployment.apps/httpd-frontend created

```