

Kubernetes have advanced networking capabilities that allow Pods and Services to communicate inside the cluster's network. An Ingress enables inbound connections to the cluster, allowing external traffic to reach the correct Pod.

Ingress enables externally-reachable urls, load balance traffic, terminate SSL, offer name based virtual hosting for a Kubernetes cluster.

In this scenario you will learn how to deploy and configure Ingress rules to manage incoming HTTP requests.

### Step 1 - Create Deployment

To start, deploy an example HTTP server that will be the target of our requests. The deployment contains three deployments, one called *webapp1* and a second called *webapp2*, and a third called *webapp3* with a service for each.

```
controlplane $ cat deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp1
spec:
  replicas: 1
  selector:
    matchLabels:
      app: webapp1
  template:
    metadata:
      labels:
        app: webapp1
    spec:
      containers:
      - name: webapp1
        image: katacoda/docker-http-server:latest
        ports:
        - containerPort: 80
```

---

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp2
spec:
  replicas: 1
  selector:
    matchLabels:
      app: webapp2
  template:
    metadata:
      labels:
        app: webapp2
    spec:
      containers:
      - name: webapp2
```

```
    image: katacoda/docker-http-server:latest
    ports:
      - containerPort: 80
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp3
spec:
  replicas: 1
  selector:
    matchLabels:
      app: webapp3
  template:
    metadata:
      labels:
        app: webapp3
    spec:
      containers:
        - name: webapp3
          image: katacoda/docker-http-server:latest
          ports:
            - containerPort: 80
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: webapp1-svc
  labels:
    app: webapp1
spec:
  ports:
    - port: 80
  selector:
    app: webapp1
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: webapp2-svc
  labels:
    app: webapp2
spec:
  ports:
    - port: 80
  selector:
    app: webapp2
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: webapp3-svc
  labels:
    app: webapp3
spec:
  ports:
```

- port: 80  
selector:  
  app: webapp3

Deploy the definitions with `kubectl apply -f deployment.yaml`

The status can be viewed with `kubectl get deployment`

```
controlplane $ kubectl apply -f deployment.yaml
deployment.apps/webapp1 created
deployment.apps/webapp2 created
deployment.apps/webapp3 created
service/webapp1-svc created
service/webapp2-svc created
service/webapp3-svc created
controlplane $ kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
webapp1       1/1     1            1           12s
webapp2       1/1     1            1           12s
webapp3       1/1     1            1           12s
controlplane $
```

## Step 2 - Deploy Ingress

The YAML file *ingress.yaml* defines a Nginx-based Ingress controller together with a service making it available on Port 80 to external connections using ExternalIPs. If the Kubernetes cluster was running on a cloud provider then it would use a LoadBalancer service type.

The ServiceAccount defines the account with a set of permissions on how to access the cluster to access the defined Ingress Rules. The default server secret is a self-signed certificate for other Nginx example SSL connections and is required by the [Nginx Default Example](#).

```
controlplane $ cat ingress.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: nginx-ingress
---
apiVersion: v1
kind: Secret
metadata:
  name: default-server-secret
  namespace: nginx-ingress
type: kubernetes.io/tls
data:
```

tls.crt:

LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUN2akNDQWFZQ0NRREFFRjI0THNhWFhEQU5CZ2txaGtpRzl3MEJBUXNGQURBaE1SOHdiUVIEVIFRRERCWk8KUjBsT1dFbHVhVW0psYzNORGIyNTBjbTlZykdWeU1CNFhEVEU0TURreE1qRTRNRE16TIZvWERUSXpNRGt4TVRFNAPnRE16TIZvd0IURWZlNQBHQTfVRUF3d1dUa2RKVGXoSmJtZHIaWE56UTI5dWRISnZiR3hsY2pDQ0FTSXdeUVIKCktvWklodmNOQVFFQkJRQURnZ0VQQUURDQ0FRb0NnZ0VCQUwvN2hIUeUeFWGRMDjNyaUM3QIBrMTNpWkt5eTlyQ08KR2xZUXYyK2EzUDF0azlrS3YwVGF5aGRcBDRrcnNUcTZzZm8vWUk1Y2VhbkW4WGM3U1pyQkVRYm9EN2REbWs1Qgo4eDZLS2xHWU5lWlG0Rm5UZOVPaStlM2ptTFFxRIBSY1kzVnNPazFFeUZBL0JnWlJVbkNHZUtGeERSN0tQdGhyCmtqSXVuektURXUyaDU4Tlp0S21ScUJHdDEwcnTRYzhZT3ExM2FnbmovUWRjc0ZYyTJnMjB1K1IYZDdoZ3krZksKWk4vVUkxQUQ0YzZyM1lma1ZWUmVHd1lxQVp1WXN2V0RKbW1GNWRwdEMzN011cDBPRUxVTEsSakZJOTZXNXlWsaO1TmdPc25NWFJNV1hYVlpiNWRxT3R0SmRtS3FhZ25TZ1JQQVpQN2MwQjFQU2FqYzZjNGZRVXpNQ0F3RUFBVEFOCKJna3Foa2lHOXcwQkFRc0ZBQU9DQVFFQWpLb2tRdGRPCesRtZhibWVPc3lySmdJSXJycVFVY2ZOuitjB0hZVUoKdGhrYnhiTFMzR3VBTWl5dm15VExPY2xeC9aYzJPblEwMEJCLzltb0swcitFZ1U2UIvrRwTWCitTFA3NTdUWgozZWl4dmdPdEduMS9ienM3bzNBaS9kclkrCUi5Q2k1S3lPc3FHTG1US2xFaUtOYkcyR1ZyTWxjS0ZYQU80YTY3CkInc1hzYktNbTQwV1U3cG9mcGltU1ZmaXFsdKv5YmN3N0NYODF6cFErUyt1eHRYK2VBZ3V0NHh3VII5d2lyVXYKelhuZk9HbWhWNTThDd1dIQnNKA0kxNXhaa2VUWXdSN0diaEFMSkZUUKk3dkhvQXprTWlzbjAxQjQyWjNnR3RXNQpJUDFmTlplOFUvOWxiUHN0T21FRFZkdjF5ZytVRVJxbStGSis2R0oxeFJGcGZnPT0KLS0tLS1FTkQgQ0VSVEIGSUNBVEU0tLS0tLQo=

tls.key:

LS0tLS1CRUdJTiBSU0EgUUFJJVkfFURSBLRVktLS0tLQpNSUJFfEFJQkFBS0NBuUUVBdi91RWM4b1JkMHUvZXXVJTHNFK1RYZUpRckxMMnNJNGFWaEMvYjYyYy9XMiRiNHVCIJOcktGMEdYaVN1eE9ycXgrajlnamx4NXFjdnhkenRKbXNFUkJ1Z1B0ME9hVGtlekhvb3FVWmcwZGxmZ1dkT0EKUTZMNTdlT1I0Q29VOUZ4amRXdzUUVVRJVUQ4R0JsRINjSVo0b1hFTkhzbysyR3VTTWk2Zk1wTVM3YUhudzFtMApxWkdvRWEzWFNyeZEJ6eGc2clhkcUNIUDICMXl3VmRyYURiUzc1aGQzdUdETDU4cGszOVFqVUFQaHpxdmRoK1JWCIZGNGJCaW9CbTVpeTIZTW1hWVhsMm0wTGZeTZuUTRRdFFZdEdNVWozcGJtdlFmazJBNNljeGRFeFpkZFZsdmwKMm82MjBsMllxcHFDZEtCRThCay90eIFIVTIKcU56cHpoOUJUTXdJREFRQUJBb0lCQVFDZkIhBxOwOHhRVmorNwplZnZlUjUXQwQ0YzR2MxNld6eDhVNml4MHg4Mm15d1kxUUNIL3BzWE9LZlRxt1h1SENYUlp5TnUvZ2lVUQ4bUFOCmxOMjRZTW0TWRJODg5TEZ0Tkp3QU5OODJDeTczckM5bzVvUDlkazAvYzRlBjAzSkVYNzZ5QjgzQm9rR1FvYksKMjhMNk0rdHUzUmFqNjd6Vmc2d2szaEhrU0pXSzBwV1YrSjdrUkRWYmhDYUZhNk5nMUZNRWxhTlozVDhhUUtyQgpDUDNDeEFTdjYxWTk5TEI4KzNXWVFIK3NYaTVGM01pYVNBZ1BkQUk3WEh1dXFET1lvMU5PL0JoSGt1aVg2QnRtCnorNTZud2pZMy8yUytSRmNBc3JMTnlwMDJZZi9oY0IraVIDNzVWYmcydVd6WTY3TWdOTGQ5VW9RU3BDRkYrVm4KM0cyUnhybnhBb0dCQU40U3M0ZVIPU2huMVpQqjdTUZsY0k2RHR2S2ErTGZTTXFyY2pOZjJlSEpZNnhubmxKdgpGenpGL2RiVWVtBwXsSekR0WkdIcXZaHfISy9iTiJlyeWJhOU1WMDIRQ0JFTk5jNmtWajJTVHpUWkJVbEx4QzYrCk93Z0wyZHhKendWelU0VC84ajdHaIRUN05BZVpFS2FvRHFyRG5BYWkyaW5oZU1JVVWZHRXFGKzJyQW9HQKFOMVAKK0tZL0lsS3RWRzRKSklQNZBjUis3RmpyeXJpY05iWCtQVzUvOXFHaWxnY2grZ3l4b25BWlBpd2NpeDN3QVpGdwpaZC96ZFB2aTBkWEppc1BSZjRMazg5b2pCUmpiRmRmc2l5UmJYbyt3TFU4NUhRU2NGMnN5aUFPaTVBRHdVU0FkCm45YWFweUNweEFkREtERHdObit3ZFhtaTZ0OHRpSFRkK3RoVDhkaVpBb0dCQUt6Wis1bG9OOTBtYlF4Vvh5YUwKMjFSUm9tMGJjcndsTmVcAWNFSmlzaEhYa2xpSVVxZ3hSZklNM2hhUVRUckIKZENFaHFsV01aV0xPb2l2NTNyZgo3aFIMSXM1ZUtka3o0aFRVdnpldm9TMHVXcm9CV2xOVHIGanlrSWhKZnZUc0hpOGdsU3FkbXgySkJhZUFVWUNXCndNdlQ4NmNLclNyNkQrZG8wS05FZzFsL0FvR0FIMkFVdHVfFbFNqLzBmRzgrV3hHc1RFV1JqcIRNUzRSUjhRWXQKeXddjFA4aDZxTGxKUTRCWGxQU05rMXZLTmtOUkxlb2pZT2pCQTViYjhibXNVU1BIV09NNENoaFJ4QnlHbmR2eAphYkJDRkFwY0lvbEg4d1R0alVZYIN5T294ZGt5OE0ek90ajJhS0FiZhd6NIaRwDZDODhjZmxYVf05MWpYL3RMCjF3TmRKS2tDZ1lCbyt0UzB5TzJ2SWFmK2UwSkN5TGhzVDQ5cTN3Zis2QWVqWGx2WDJ1VnRYejN5QTZnbXo5aCsKcDNlK2JMRUxwb3B0WFhNdUFRR0xhUkcrYINNCjR5dERYbE5ZSndUeThXczNKY3dlSTdqZVp2b0ZpbmNvVlVIMwphdmxoTUVCRGYxSjltSDB5cDBwWUNaS2ROdHNvZEZtQktzVEtQMjJhTmtsVvhCS3gyZzR6cFE9PQotLS0tLUVORCBSU0EgUUFJJVkfFURSBLRVktLS0tLQo=

---

apiVersion: v1

kind: ServiceAccount

metadata:

name: nginx-ingress

namespace: nginx-ingress

```

---
kind: ConfigMap
apiVersion: v1
metadata:
  name: nginx-config
  namespace: nginx-ingress
data:
---
# Described at: https://docs.nginx.com/nginx-ingress-controller/installation/installation-with-manifests/
# Source from:
https://github.com/nginxinc/kubernetes-ingress/blob/master/deployments/common/ingress-class.yaml
apiVersion: networking.k8s.io/v1beta1
kind: IngressClass
metadata:
  name: nginx
  # annotations:
  # ingressclass.kubernetes.io/is-default-class: "true"
spec:
  controller: nginx.org/ingress-controller
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-ingress
  namespace: nginx-ingress
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx-ingress
  template:
    metadata:
      labels:
        app: nginx-ingress
    spec:
      serviceAccountName: nginx-ingress
      containers:
        - image: nginx/nginx-ingress:edge
          imagePullPolicy: Always
          name: nginx-ingress
          ports:
            - name: http
              containerPort: 80
            - name: https
              containerPort: 443
      env:
        - name: POD_NAMESPACE
          valueFrom:
            fieldRef:
              fieldPath: metadata.namespace
        - name: POD_NAME
          valueFrom:
            fieldRef:
              fieldPath: metadata.name
      args:
        - --nginx-configmaps=$(POD_NAMESPACE)/nginx-config

```

```

- -default-server-tls-secret=$(POD_NAMESPACE)/default-server-secret
---
apiVersion: v1
kind: Service
metadata:
  name: nginx-ingress
  namespace: nginx-ingress
spec:
  type: NodePort
  ports:
    - port: 80
      targetPort: 80
      protocol: TCP
      name: http
    - port: 443
      targetPort: 443
      protocol: TCP
      name: https
  selector:
    app: nginx-ingress
  externalIPs:
    - 10.0.0.6

```

## Task

The Ingress controllers are deployed in a familiar fashion to other Kubernetes objects with `kubectl create -f ingress.yaml`

The status can be identified using `kubectl get deployment -n nginx-ingress`

```

controlplane $ kubectl create -f ingress.yaml
namespace/nginx-ingress created
secret/default-server-secret created
serviceaccount/nginx-ingress created
configmap/nginx-config created
ingressclass.networking.k8s.io/nginx created
deployment.apps/nginx-ingress created
service/nginx-ingress created
controlplane $ kubectl get deployment -n nginx-ingress
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-ingress        0/1     1             0           3s

```

## Step 3 - Deploy Ingress Rules

Ingress rules are an object type with Kubernetes. The rules can be based on a request host (domain), or the path of the request, or a combination of both.

An example set of rules are defined within `cat ingress-rules.yaml`

The important parts of the rules are defined below.

The rules apply to requests for the host *my.kubernetes.example*. Two rules are defined based on the path request with a single catch all definition. Requests to the path */webapp1* are forwarded onto the service *webapp1-svc*. Likewise, the requests to */webapp2* are forwarded to *webapp2-svc*. If no rules apply, *webapp3-svc* will be used.

This demonstrates how an application's URL structure can behave independently about how the applications are deployed.

```
controlplane $ cat ingress-rules.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: webapp-ingress
spec:
  ingressClassName: nginx
  rules:
  - host: my.kubernetes.example
    http:
      paths:
      - path: /webapp1
        backend:
          serviceName: webapp1-svc
          servicePort: 80
      - path: /webapp2
        backend:
          serviceName: webapp2-svc
          servicePort: 80
      - backend:
          serviceName: webapp3-svc
          servicePort: 80
```

### Task

As with all Kubernetes objects, they can be deployed via

```
kubectl create -f ingress-rules.yaml
```

Once deployed, the status of all the Ingress rules can be discovered via

```
kubectl get ing
```

```
controlplane $ kubectl create -f ingress-rules.yaml
ingress.extensions/webapp-ingress created
controlplane $ kubectl get ing
NAME                CLASS    HOSTS                ADDRESS    PORTS    AGE
webapp-ingress      nginx    my.kubernetes.example 80         4s
```

#### Step 4 - Test

With the Ingress rules applied, the traffic will be routed to the defined place.

The first request will be processed by the *webapp1* deployment.

```
curl -H "Host: my.kubernetes.example" 10.0.0.6/webapp1
```

The second request will be processed by the *webapp2* deployment.

```
curl -H "Host: my.kubernetes.example" 10.0.0.6/webapp2
```

Finally, all other requests will be processed by *webapp3* deployment.

```
curl -H "Host: my.kubernetes.example" 10.0.0.6
```