

A pod definition file nginx.yaml is given. Create a pod using the file.
Only create the POD for now. We will inspect its status next.
Use the command `kubectl create -f nginx.yaml`

```
controlplane ~ → kubectl create -f nginx.yaml
pod/nginx created

controlplane ~ → cat nginx.yaml
---
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - image: nginx
    name: nginx
```

What is the status of the created POD?
Run the command: `kubectl get pods` and check the Status column.

```
controlplane ~ → kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     0/1     Pending   0           39s
```

Why is the POD in a pending state? No scheduler present
Inspect the environment for various kubernetes control plane components.
Run the command: `kubectl get pods --namespace kube-system` to see the status of scheduler pod. We have removed the scheduler from this Kubernetes cluster. As a result, as it stands, the pod will remain in a pending state forever.

```
controlplane ~ → kubectl get pods --namespace kube-system
NAME                                READY   STATUS    RESTARTS   AGE
coredns-787d4945fb-dzg2m           1/1     Running   0           6m22s
coredns-787d4945fb-vvt9g           1/1     Running   0           6m22s
etcd-controlplane                   1/1     Running   0           6m32s
kube-apiserver-controlplane          1/1     Running   0           6m37s
kube-controller-manager-controlplane 1/1     Running   0           6m36s
kube-proxy-8fq79                    1/1     Running   0           6m22s
kube-proxy-kxz4j                     1/1     Running   0           6m9s
```

Manually schedule the pod on node01.
Delete and recreate the POD if necessary.

```
controlplane ~ → kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     0/1     Pending   0           3m15s

controlplane ~ → kubectl delete pods nginx
pod "nginx" deleted

controlplane ~ → kubectl get nodes
NAME          STATUS    ROLES          AGE      VERSION
controlplane  Ready     control-plane   8m30s    v1.26.0
node01        Ready     <none>          7m59s    v1.26.0

controlplane ~ → vi nginx.yaml

controlplane ~ → cat nginx.yaml
---
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  nodeName: node01
  containers:
  - image: nginx
    name: nginx

controlplane ~ → kubectl create -f nginx.yaml
pod/nginx created

controlplane ~ → kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE    NOMINATED NODE   READINESS GATES
nginx     1/1     Running   0           8s    10.244.1.2   node01   <none>            <none>
```

Now schedule the same pod on the controlplane node.
Delete and recreate the POD if necessary.

```
controlplane ~ → kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     0/1     Pending   0           3m15s

controlplane ~ → kubectl delete pods nginx
pod "nginx" deleted

controlplane ~ → kubectl get nodes
NAME          STATUS    ROLES          AGE      VERSION
controlplane  Ready     control-plane   8m30s    v1.26.0
node01        Ready     <none>          7m59s    v1.26.0
```

```
controlplane ~ → cat nginx.yaml
---
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  nodeName: controlplane
  containers:
  - image: nginx
    name: nginx

controlplane ~ → kubectl create -f nginx.yaml
pod/nginx created

controlplane ~ → kubectl get pods -o wide
```

| NAME | READY | STATUS | RESTARTS | AGE | IP | NODE | NOMINATED NODE | READINESS GATES |
|-------|-------|---------|----------|-----|------------|--------------|----------------|-----------------|
| nginx | 1/1 | Running | 0 | 8s | 10.244.0.4 | controlplane | <none> | <none> |

Labels and selectors :

We have deployed a number of PODs. They are labelled with tier, env and bu. How many PODs exist in the dev environment (env)?

Use selectors to filter the output

Run the command `kubectl get pods --selector env=dev --no-headers | wc -l`

```
controlplane ~ → kubectl get pods
```

| NAME | READY | STATUS | RESTARTS | AGE |
|-------------|-------|---------|----------|-----|
| db-1-7nmtv | 1/1 | Running | 0 | 52s |
| db-1-tfc79 | 1/1 | Running | 0 | 51s |
| app-1-4nw2w | 1/1 | Running | 0 | 52s |
| app-2-dn1r5 | 1/1 | Running | 0 | 52s |
| auth | 1/1 | Running | 0 | 52s |
| app-1-f75sd | 1/1 | Running | 0 | 52s |
| app-1-qmbx8 | 1/1 | Running | 0 | 52s |
| app-1-zzxdf | 1/1 | Running | 0 | 51s |
| db-1-c7ttr | 1/1 | Running | 0 | 52s |
| db-2-zptnj | 1/1 | Running | 0 | 51s |
| db-1-vjdzb | 1/1 | Running | 0 | 52s |

```
controlplane ~ → kubectl get pods --selector env=dev --no-headers
```

| | | | | |
|-------------|-----|---------|---|-----|
| db-1-7nmtv | 1/1 | Running | 0 | 95s |
| db-1-tfc79 | 1/1 | Running | 0 | 94s |
| app-1-4nw2w | 1/1 | Running | 0 | 95s |
| app-1-f75sd | 1/1 | Running | 0 | 95s |
| app-1-qmbx8 | 1/1 | Running | 0 | 95s |
| db-1-c7ttr | 1/1 | Running | 0 | 95s |
| db-1-vjdzb | 1/1 | Running | 0 | 95s |

```
controlplane ~ → kubectl get pods --selector env=dev --no-headers | wc -l
7
```

How many PODs are in the finance business unit (bu)?

Run the command `kubectl get pods --selector bu=finance --no-headers | wc -l`

```
controlplane ~ → kubectl get pods --selector bu=finance --no-headers | wc -l
6
```

How many objects are in the prod environment including PODs, ReplicaSets and any other objects?

Run the command to get exact number of objects `kubectl get all --selector env=prod --no-headers | wc -l`

```
controlplane ~ → kubectl get all --selector env=prod --no-headers | wc -l
7
```

Identify the POD which is part of the prod environment, the finance BU and of frontend tier?

Run the command `kubectl get all --selector env=prod,bu=finance,tier=frontend`

```
controlplane ~ → kubectl get all --selector env=prod,bu=finance,tier=frontend
NAME                READY   STATUS    RESTARTS   AGE
pod/app-1-zzxdf     1/1     Running   0           4m57s
```

A ReplicaSet definition file is given `replicaset-definition-1.yaml`. Try to create the replicaset. There is an issue with the file. Try to fix it.

```
controlplane ~ → cat replicaset-definition-1.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: replicaset-1
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: front-end
  template:
    metadata:
      labels:
        tier: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```

```
controlplane ~ → cat replicaset-definition-1.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: replicaset-1
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: front-end
  template:
    metadata:
      labels:
        tier: front-end
    spec:
      containers:
        - name: nginx
          image: nginx

controlplane ~ → kubectl apply -f replicaset-definition-1.yaml
replicaset.apps/replicaset-1 created
```

Taints and Tolerations

```
controlplane ~ → kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
controlplane        Ready     control-plane  2m24s  v1.26.0
node01              Ready     <none>      111s   v1.26.0

controlplane ~ → kubectl get nodes | wc -l
3

controlplane ~ → kubectl get nodes --no-headers | wc -l
2
```

Do any taints exist on node01 node?

Run the command: `kubectl describe node node01 | grep -i taints` to check taint exists

```
controlplane ~ → kubectl describe node node01 | grep -i taints
Taints:
<none>

controlplane ~ → kubectl describe node node01
Name: node01
Roles: <none>
Labels: beta.kubernetes.io/arch=amd64
        beta.kubernetes.io/os=linux
        kubernetes.io/arch=amd64
        kubernetes.io/hostname=node01
        kubernetes.io/os=linux
Annotations: flannel.alpha.coreos.com/backend-data: {"VNI":1,"VtepMAC":"ca:5a:41:14:22:f8"}
             flannel.alpha.coreos.com/backend-type: vxlan
             flannel.alpha.coreos.com/kube-subnet-manager: true
             flannel.alpha.coreos.com/public-ip: 172.25.0.6
             kubeadm.alpha.kubernetes.io/cni-socket: unix:///var/run/containerd/containerd.sock
             node.alpha.kubernetes.io/ttl: 0
             volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Tue, 25 Apr 2023 03:17:40 -0400
Taints: <none>
Unschedulable: false
Lease:
  HolderIdentity: node01
  AcquireTime: <unset>
  RenewTime: Tue, 25 Apr 2023 03:20:54 -0400
Conditions:
  Type           Status  LastHeartbeatTime           LastTransitionTime           Reason
  ----           -
NetworkUnavailable False   Tue, 25 Apr 2023 03:17:49 -0400 Tue, 25 Apr 2023 03:17:49 -0400 Flann
MemoryPressure  False   Tue, 25 Apr 2023 03:18:10 -0400 Tue, 25 Apr 2023 03:17:40 -0400 Kube
DiskPressure    False   Tue, 25 Apr 2023 03:18:10 -0400 Tue, 25 Apr 2023 03:17:40 -0400 Kube
```

Create a taint on node01 with key of spray, value of mortein and effect of NoSchedule

Run the command: `kubectl taint nodes node01 spray=mortein:NoSchedule`

```
controlplane ~ → kubectl taint nodes node01 spray=mortein:NoSchedule
node/node01 tainted
```

Create a new pod with the nginx image and pod name as mosquito.

```
controlplane ~ → kubectl run mosquito --image=nginx
pod/mosquito created
```

Solution manifest file to create a pod called mosquito as follows:

```
---
apiVersion: v1
kind: Pod
metadata:
  name: mosquito
spec:
  containers:
  - image: nginx
    name: mosquito
```

then run `kubectl create -f <FILE-NAME>.yaml`

What is the state of the POD?

```
controlplane ~ → kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
mosquito      0/1     Pending   0           68s
```

Why do you think the pod is in a pending state? Pod mosquito can't tolerate taint mortein.

```
node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason             Age   From          Message
  ----      -
Warning    FailedScheduling   107s  default-scheduler  0/2 nodes are available: 1 node(s) had untolerated taint {node-role.kubernetes.io/control-plane: }, 1 node(s) had untolerated taint {spray: mortein}. Preemption: 0/2 nodes are available: 2 Preemption is not helpful for scheduling..
```

Create another pod named bee with the nginx image, which has a toleration set to the taint mortein.

```
controlplane ~ → cat nee.yaml
---
apiVersion: v1
kind: Pod
metadata:
  name: bee
spec:
  containers:
  - image: nginx
    name: bee
  tolerations:
  - key: spray
    value: mortein
    effect: NoSchedule
    operator: Equal

controlplane ~ → kubectl create -f nee.yaml
pod/bee created
```

Notice the bee pod was scheduled on node node01 despite the taint.

```
controlplane ~ → kubectl create -f nee.yaml
pod/bee created

controlplane ~ → kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE    NOMINATED NODE   READINESS GATES
bee       1/1     Running   0           35s   10.244.1.2   node01   <none>            <none>
mosquito  0/1     Pending   0           4m29s <none>      <none>   <none>            <none>
```

Do you see any taints on controlplane node?

Run the command: `kubectl describe node controlplane` and see the taint property.

```
controlplane ~ → kubectl describe node controlplane | grep -i taint
Taints:          node-role.kubernetes.io/control-plane:NoSchedule
```

Remove the taint on controlplane, which currently has the taint effect of NoSchedule.

Run the command: `kubectl taint nodes controlplane node-role.kubernetes.io/control-plane:NoSchedule-` to untaint the node.

```
controlplane ~ → kubectl taint nodes controlplane node-role.kubernetes.io/control-plane:NoSchedule-
error: node controlplane already has node-role.kubernetes.io/control-plane taint(s) with same effect(s) and --overwrite is false

controlplane ~ ✗ kubectl taint nodes controlplane node-role.kubernetes.io/control-plane:NoSchedule-
node/controlplane untainted
```


What is the state of the pod mosquito now?

```
controlplane ~ → kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE       NOMINATED NODE   READINESS GATES
bee       1/1     Running   0           3m37s  10.244.1.2   node01     <none>           <none>
mosquito  1/1     Running   0           7m31s  10.244.0.4   controlplane <none>           <none>
```

Which node is the POD mosquito on now?

Controlplane

NodeAffinity

How many Labels exist on node node01?

Run the command `kubectl describe node node01` and count the number of labels.

```
controlplane ~ → kubectl describe node node01
Name:          node01
Roles:         <none>
Labels:        beta.kubernetes.io/arch=amd64
               beta.kubernetes.io/os=linux
               kubernetes.io/arch=amd64
               kubernetes.io/hostname=node01
               kubernetes.io/os=linux
Annotations:   flannel.alpha.coreos.com/backend-data: {"VNI":1,"VtepMAC":"c2:3a:7f:a8:e4:b4"}
               flannel.alpha.coreos.com/backend-type: vxlan
               flannel.alpha.coreos.com/kube-subnet-manager: true
               flannel.alpha.coreos.com/public-ip: 172.25.0.80
               kubeadm.alpha.kubernetes.io/cri-socket: unix:///var/run/containerd/containerd.sock
               node.alpha.kubernetes.io/ttl: 0
               volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Tue, 25 Apr 2023 03:54:30 -0400
Taints:         <none>
Unschedulable:  false
Lease:          <none>
```

What is the value set to the label key `beta.kubernetes.io/arch` on node01?

Apply a label `color=blue` to node node01

Run the command: `kubectl label node node01 color=blue`

```
controlplane ~ → kubectl label node node01 color=blue
node/node01 labeled
```

Create a new deployment named blue with the nginx image and 3 replicas.

Run the command: `kubectl create deployment blue --image=nginx --replicas=3`

```
controlplane ~ → kubectl create deployment blue --image=nginx --replicas=3
deployment.apps/blue created
```

Which nodes can the pods for the blue deployment be placed on? Controlplane and node01

Make sure to check taints on both nodes!

Check if controlplane and node01 have any taints on them that will prevent the pods to be scheduled on them.

If there are no taints, the pods can be scheduled on either node.

So run the following command to check the taints on both nodes.

`kubectl describe node controlplane | grep -i taints`

`kubectl describe node node01 | grep -i taints`

```
controlplane ~ → kubectl describe node controlplane | grep -i taints
Taints:         <none>

controlplane ~ → kubectl describe node node01 | grep -i taints
Taints:         <none>
```


Set Node Affinity to the deployment to place the pods on node01 only.
Edit the deployment blue and add the Node Affinity with specified key and value.

Update the deployment by running `kubectl edit deployment blue` and add the `nodeaffinity` section as follows:

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: blue
spec:
  replicas: 3
  selector:
    matchLabels:
      run: nginx
  template:
    metadata:
      labels:
        run: nginx
    spec:
      containers:
      - image: nginx
        imagePullPolicy: Always
        name: nginx
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
            - matchExpressions:
              - key: color
                operator: In
                values:
                - blue
```

```
controlplane ~ → kubectl edit deployment blue
deployment.apps/blue edited
```

Which nodes are the pods placed on now?

```
controlplane ~ → kubectl get pods -o wide
```

| NAME | READY | STATUS | RESTARTS | AGE | IP | NODE | NOMINATED | NODE | READINESS | GATES |
|-----------------------|-------|---------|----------|-----|------------|--------|-----------|------|-----------|-------|
| blue-7cf59b987f-6d7wv | 1/1 | Running | 0 | 57s | 10.244.1.6 | node01 | <none> | | <none> | |
| blue-7cf59b987f-6qxfs | 1/1 | Running | 0 | 65s | 10.244.1.5 | node01 | <none> | | <none> | |
| blue-7cf59b987f-fdd7v | 1/1 | Running | 0 | 68s | 10.244.1.4 | node01 | <none> | | <none> | |

Create a new deployment named red with the nginx image and 2 replicas, and ensure it gets placed on the controlplane node only.

Use the label key - `node-role.kubernetes.io/control-plane` - which is already set on the controlplane node.

```
controlplane ~ → cat red.yaml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: red
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      run: nginx
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        run: nginx
```

```
    spec:
```

```
      containers:
```

```
      - image: nginx
```

```
        imagePullPolicy: Always
```

```
        name: nginx
```

```
      affinity:
```

```
        nodeAffinity:
```

```
          requiredDuringSchedulingIgnoredDuringExecution:
```

```
            nodeSelectorTerms:
```

```
              - matchExpressions:
```

```
                - key: node-role.kubernetes.io/control-plane
```

```
                  operator: Exists
```

```
controlplane ~ → kubectl create -f red.yaml
```

```
deployment.apps/red created
```